



Algorithms

Network Flow



Hee-Kap Ahn
Graduate School of Artificial Intelligence
Dept. Computer Science and Engineering
Pohang University of Science and Technology (POSTECH)

Flows in Networks



We are given

- a **directed graph** $G = (V, E)$;
- two special nodes $s, t \in V$, a **source** and a **sink** of G , respectively; and
- **capacities** $c_e > 0$.

We want to **send as much flow as possible** from s to t such that

- $0 \leq f_e \leq c_e$ for all $e \in E$.
- flow is conserved : $\sum_{(w,u) \in E} f_{wu} = \sum_{(u,z) \in E} f_{uz}$.
- $\text{size}(f) = \sum_{(s,u) \in E} f_{su} = \sum_{(v,t) \in E} f_{vt} = \text{size}(f)$.

The maximum flow problem reduces to linear programming!

Residual Networks


The general simplex algorithm would

- start with zero flow
- repeat:
choose an appropriate path from s to t , and increase flow along the edges of the path as much as possible.


Each iteration of simplex looks for an $s - t$ path whose edges (v, w) can be of two types:

- $(v, w) \in E$ and $f_{vw} < c_{vw}$.
- $(w, v) \in E$ and $f_{wv} > 0$.


$(a, b) \in E$ and $f_{ab} = 1$. Path $sbat$ is chosen.



Given network




Final flow



Input graph $G = (V, E)$ vs. Residual graph $G^f = (V, E^f)$.


For an edge $(v, w) \in E$ with $f_{vw} \leq c_{vw}$, $G^f = (V, E^f)$ has two edges $(v, w), (w, v)$ with residual capacities

$$\begin{aligned} c_{vw}^f &= c_{vw} - f_{vw}, \\ c_{wv}^f &= f_{vw}. \end{aligned}$$




Simplex Algorithm in Residual Networks

Current flow




Residual graph




A $s - t$ path *scbde* of flow 1.

Simplex Algorithm in Residual Networks

Current flow




Residual graph




A $s - t$ path $sabt$ of flow 1.

Simplex Algorithm in Residual Networks

Current flow




Residual graph




A $s - t$ path *sdet* of flow 4.

Simplex Algorithm in Residual Networks

Current flow




Residual graph




A $s - t$ path $sacbt$ of flow 1.

Simplex Algorithm in Residual Networks

Current flow




Residual graph




The final flow and the Residual graph.

Simplex Algorithm in Residual Networks

Current flow



Residual graph




The final flow and the Residual graph.


Flow and (s, t) -cut

A certificate of optimality Partition the nodes of the network into two groups. In the following residual graph, let $L = \{s, a, c\}$ and $R = \{b, d, e, t\}$, for example.

Residual graph G^f



(L, R) of G



An (s, t) -cut partitions the vertices into two disjoint groups L and R s.t. $s \in L$ and $t \in R$. Its capacity is the total capacity of edges $\in G$ **from** L **to** R .

Pick any flow f and any (s, t) -cut.
Then $\text{size}(f) \leq \text{capacity}(L, R)$.


Max-flow Min-cut Theorem

The size of the maximum flow in a network equals the capacity of the smallest (s, t) -cut.

Suppose f is the final flow when the algorithm terminates. Then t is no longer reachable from s in G^f .

Let L be the set of nodes reachable from s in G^f , and $R = V - L$.

We claim that $\text{size}(f) = \text{capacity}(L, R)$.




Consider any edge $e \in E$ from a node in L to a node in R . Then $f_e = c_e$. Consider now any edge $e' \in E$ from a node in R to a node in L . Then $f_{e'} = 0$. Do you see why?

Therefore, (L, R) is the smallest (s, t) -cut!

Flows in Networks

Efficiency Each iteration of our maximum-flow algorithm is efficient, requiring $O(|E|)$ time, if a DFS or BFS is used to find an $s - t$ path. But how many iterations are there?




Is there a way to avoid $C|E|$ iterations, when C is a huge number?

Bipartite Matching

We have a set of people P and a set of jobs J .
Each person can do some of the jobs.
We can model this as a bipartite graph.

A **matching** gives an assignment of people to jobs such that


- a person is assigned to at most one job and
- at most one person is assigned to a job.



Bipartite Matching

Given an instance of bipartite matching, we can find a maximum matching by the following reduction.


- Replace each edge to a directed edge from a person to a job of capacity 1.
- Add a special node s and connect s to every node in People by directed edge of capacity 1.
- Add a special node t and connect every node in Jobs to t by directed edge of capacity 1.



Bipartite Matching

Given an instance of bipartite matching, we can find a maximum matching by the following reduction.


- Replace each edge to a directed edge from a person to a job of capacity 1.
- Add a special node s and connect s to every node in People by directed edge of capacity 1.
- Add a special node t and connect every node in Jobs to t by directed edge of capacity 1.



Bipartite Matching

Given an instance of bipartite matching, we can find a maximum matching by the following reduction.

- Replace each edge to a directed edge from a person to a job of capacity 1.
- Add a special node s and connect s to every node in People by directed edge of capacity 1.
- Add a special node t and connect every node in Jobs to t by directed edge of capacity 1.



Bipartite Matching

The edges used in the maximum flow correspond to the maximum matching.

- all capacities are integers, so the maximum flow is integral.
- all capacities are 1, so an edge is used (flow 1) or not.

Let M be the set of edges used in the maximum flow of value k . Then M is a matching because there is at most one edge in M leaving any person and entering any job. Moreover, M consists of k edges.

If there is a matching consisting of more than k edges, there must be a flow of value larger than k , which is a contradiction.

For n people and m edges, the running time is?

